Signal: Agent-Powered Growth and Marketing

Product conclusion

Signal should operate as a founder's growth workspace: identify a credible buyer problem, propose a small experiment, produce useful materials, support deliberate distribution, and connect the work to observed customer outcomes. The economic objective is paying customers and retained value. Model activity, draft counts, impressions, and generated images are intermediate outputs.

The appropriate starting strategy for Lemonade and CineGen is a small portfolio of zero-cash tests. A broad channel library prevents familiar-channel bias, while a short active list protects the founder's time. The app should support every major traction channel without suggesting that a solo founder execute all of them simultaneously. Channel choice must reflect the business, buyer, buying process, and evidence, rather than the model's confidence alone. Published practitioner guidance similarly distinguishes acquisition approaches by business type and audience. [1]

The resulting design combines the existing Bullseye experiment engine with a campaign production workspace, customer conversation pipeline, measured outcomes, optional GPT execution, and local image/video tooling. Automation prepares and organizes real work. It does not turn an unknown market into a guaranteed revenue stream. First-customer stories are useful qualitative examples, not controlled evidence that one tactic will succeed for every product. [2]

Evidence and interpretation

This report uses official technical documentation, model publishers, product websites, and firsthand growth practitioners. Technical sources establish API contracts and supported workflows. Product websites establish what the businesses currently say about themselves; they do not independently validate every advertised capability. Growth recommendations and test designs are analytical judgments to validate with customer behavior.

Sources were checked during September 13, 2026. Platform access, model releases, commercial terms, and product offers can change. Each in-app channel playbook links to the supporting source context. The suggested cohort sizes, schedules, and stop rules are operational defaults authored for Signal, not universal statistical findings.

No defensible finite catalogue contains literally every way to market a business. The practical coverage target is the 19 broad traction channels, supported by reusable assets, research, customer conversations, experiments, and measurement. New tactics fit into those channels and can be expressed as campaign briefs. This is more durable than hardcoding a directory of fashionable tools.

Business positioning and the first offer

Lemonade's current website describes a private desktop assistant for Windows 11 and Linux, with local file work, screen capture, document creation, and managed Ollama/model setup. It advertises a $49 perpetual license and a 14-day trial. Its privacy explanation distinguishes local inference from network activity for setup, licensing, and optional approved tools.[3] Signal's older configuration incorrectly said an existing Ollama installation was required. That stale statement was corrected, along with an overly absolute network claim.

For Lemonade, the proposed initial experiment is a concrete document workflow for a narrow privacy-motivated audience. A strong demonstration shows a real input, the actual application, the resulting editable document, and the limitations encountered. A generic assertion about replacing every cloud assistant is a weaker test because it conflates many buying motivations. The initial call to action should invite the buyer to try the same bounded workflow and describe friction.

CineGen's public page presents a browser video-production workspace spanning storyboards, clips, timeline editing, and review. Its free offer includes one project, watermarked 720p export, and introductory AI tokens; the 3D scene creator is described as beta.[4] Signal should preserve that distinction when creating promotional copy. A proposed first test is a short storyboard-to-export walkthrough for one creator use case, with the final video and the actual timeline both visible.

For both businesses, the brand profile records audience, painful task, offer, proof, voice, constraints, and visual direction. These fields are operational context for every campaign producer. They should state what is known, which claims require a demonstration, and what must not be promised. A model should not infer unsupported pricing or capability from a product name.

Channel coverage and execution

The following table is an execution map, not a ranking of expected returns. The default first test is intentionally small. Its purpose is to expose buyer interest and practical distribution constraints before a larger commitment.

Channel	Initial test	Materials Signal prepares	Outcome to inspect
Targeting blogs	Offer an original tutorial to three relevant editors	Guest article, editor pitch, evidence notes	Qualified readers progressing to a trial
Publicity	Pitch a verifiable story to beat-matched journalists	Short pitch, fact sheet, source list	Relevant coverage and inbound conversations
Unconventional PR	Demonstrate a surprising, truthful product use	Demo script, visual brief, native posts	Attention from actual prospective buyers
Search advertising	Research one high-intent keyword cluster before funding	Landing copy and ad angles	Cost per activated lead and paid customer

Social/display advertising	Try several angles organically before paid distribution	Posts, image briefs, short scripts	Qualified conversion after the click
Offline advertising	Validate a suitable placement before printing	Concise creative and tracked destination	Attributable conversations and purchases
SEO	Answer one recurring buyer question with original detail	Complete article and supporting visual brief	Relevant organic visits and activation
Content marketing	Teach one workflow and adapt it for two destinations	Tutorial, email adaptation, social version	Engaged prospective buyers and trials
Email marketing	Draft an activation sequence for an opted-in cohort	Welcome, useful example, obstacle-removal copy	Activation, conversions, opt-outs
Viral marketing	Test an optional shareable outcome with satisfied users	Invitation and recipient landing copy	Invited people who activate and return
Engineering as marketing	Manually deliver a proposed free tool's result	Tool brief, landing copy, demonstration	Qualified free-tool users entering paid workflow
Partnerships	Propose one joint tutorial to complementary businesses	Partner proposal, event outline, follow-up	Partner-attributed activation and sales
Sales	Have five specific problem conversations	Contextual invitation and discovery questions	Qualified conversation to paid pilot
Affiliates	Prepare a small partner pilot with owner-defined terms	Partner kit, tracked offer, program outline	Net contribution after commissions/refunds
Existing platforms	Prepare one complete relevant listing	Listing copy, genuine screenshots, onboarding steps	Listing-to-activation conversion
Trade shows	Arrange buyer meetings before buying a booth	Meeting request, demonstration, follow-up	Realized sales against full participation cost
Events	Run a small practical workshop	Agenda, invitation, exercise, useful takeaway	Relevant attendance and voluntary follow-up
Speaking	Offer a practical talk to relevant hosts	Talk proposal, script, resource offer	Qualified conversations after the session
Community building	Help people with a recurring problem	Useful answers and recurring educational material	Returning participants and product adoption

Each in-app playbook includes its suitable audience, first test, sequential actions, suggested materials, measurement, reconsideration rule, and links. Campaign briefs can choose any channel and combine up to six formats. The existing Growth lab keeps the broader research, shortlist, test goals, and evidence-based focus decision.

Customer conversations deserve particular attention at zero revenue. Their value is not only sales conversion: they reveal language, alternatives, urgency, and reasons not to buy. Signal's lead pipeline therefore records the observed problem, source, stage, contact permission, notes, and follow-up date. It does not mistake a scraped name for a qualified lead or a drafted message for an actual conversation.

Product-led acquisition is a separate design possibility. A useful shareable outcome, collaboration invitation, or free adjacent tool may naturally distribute the product, but the recipient must receive independent value. Practitioner discussion of product-led acquisition supports considering product mechanisms alongside purchased distribution.[5] In Signal, these ideas become measurable campaigns or engineering requests, rather than claims that a referral loop already exists.

Research and content quality

Good automation begins with evidence. Before writing, the agent should inspect current product facts, relevant stored observations, and actual customer questions. Public web content and tool results remain reference data, not instructions that can redirect tool execution. A stored model hypothesis must retain its status as a hypothesis when recalled later.

Search content should provide a complete, original answer. Google's guidance emphasizes useful information, demonstrated experience, clear sourcing, and a satisfying answer for an intended audience.[6] Its guidance on generative AI warns against producing many low-value pages primarily to manipulate rankings.[7] The recommended workflow is one useful piece grounded in experience, followed by deliberate adaptation and measurement.

Campaign quality is evaluated in four layers. First, the material must satisfy the output contract and contain finished copy. Second, it must match the audience, offer, channel, and intended action. Third, claims and evidence must survive owner review. Fourth, actual response and conversion data determine whether the creative was useful. A structurally valid JSON response is necessary for automation but says little about persuasion or truth.

Signal's campaign producer writes a title, body, rationale, evidence references, and review notes for each requested format. It flags placeholders, numerical claims, guarantees, and testimonial language for attention. Approval requires the owner to check claims, voice, and destination. Editing returns material to draft status. Published copy is preserved as history rather than silently rewritten.

Small local models need especially clear boundaries. In live testing, an unconstrained JSON-mode response placed large portions of a social post in unexpected object keys. A closed structured-output schema corrected that concrete failure. The broader lesson is to constrain shape, validate semantics, and test real models with the actual task context. The application must not hide failed production behind a success message.

Local agents and GPT

Signal owns the agent orchestration: durable tasks, scheduling, tool dispatch, application validation, cancellation, retry, and artifacts. The inference engines are interchangeable infrastructure. The local option uses installed Ollama or a compatible loopback server. The optional GPT path uses a configured OpenAI API key and a fixed OpenAI HTTPS destination.

Ollama's native chat API supports JSON or schema-constrained formatting, model options, and model unloading through keep-alive configuration.[8] Its download API supplies model installation progress.[9] Signal uses a 16K context for its larger operator and campaign tasks, manages one local task at a time, and can start an installed Ollama server. Managed launches disable Ollama's cloud features and limit simultaneous loaded models.[10]

The desktop settings screen separates provider selection, cloud model ID, API key, output-token allowance, daily request cap, local endpoint, and local task models. No GPT requests occur merely because a key exists: the selected provider controls routing. The daily cap counts attempted generation calls conservatively; it is not a dollar cap because model prices and token usage vary. The app should not invent a cost estimate from a hardcoded price table.

For GPT, campaign output can use a closed JSON schema. The operator's heterogeneous tool input is more dynamic, so its response uses JSON mode followed by application-level validation. OpenAI documents the distinction between JSON validity and schema compliance and requires closed objects for strict structured output.[11] Both transports still pass through the same application checks.

The API key is stored in the main process using Electron's OS encryption facility and is never returned to the renderer after saving. Electron documents platform-specific protection and a Linux fallback that does not provide an adequate secret store; Signal refuses that fallback for persisted keys.[12] An environment variable remains available for CLI use or systems without encrypted storage. Remote error bodies are not shown, because they may include sensitive request context.

Creative production architecture

A strong marketing asset combines a clear idea, accurate product evidence, coherent visual direction, legible copy, and suitable distribution. Image generation alone does not provide that combination. The recommended process is to write the campaign angle, capture actual product behavior, generate illustrative supporting imagery where useful, compose the final layout, and review it at the intended display size.

ComfyUI is a suitable local rendering backend because it exposes workflow submission, queue inspection, history, and output retrieval through documented HTTP routes.[13] Signal persists a workflow reference, substituted prompt, workflow hash, job ID, endpoint, status, and output descriptors. A completed render must contain a saved output before the app reports success. A lost acknowledgement requires queue reconciliation to avoid duplicate rendering.

Three ready-to-inspect recipes are included. FLUX.1 Schnell supplies an image path with a straightforward permissive model license. Krea 2 Turbo provides an alternative already installed on the tested machine. Wan 2.2 5B supplies a short-video path. Recipe readiness checks compare required nodes and model filenames with the connected ComfyUI server. Missing components remain visible and block submission.

Option	Intended use	Resource and setup considerations	Evidence
FLUX.1 Schnell	Fast image concepts and supporting visuals	Separate diffusion, encoder, and VAE files; offloading/quantization affect latency	Official model card and ComfyUI workflow [14] [15]
Krea 2 Turbo	Editorial and stylized image exploration	Large diffusion and text-encoder weights; installed 10 GB GPU requires offloading	Official workflow and community license [16] [17]
Wan 2.2 TI2V 5B	Short illustrative text-to-video clips	Native ComfyUI offloading is documented for 8 GB VRAM; test clip length and resolution locally	Official ComfyUI guide and model card [18] [19]
LTX family	Future synchronized audio/video workflow exploration	Different model/node requirements; do not assume compatibility with a Wan workflow	Official LTX tutorial [20]

The hardware probe for this workspace found an NVIDIA RTX 3080 with 10 GB VRAM. An existing portable ComfyUI installation contained Krea 2 Turbo FP8, its text encoder, and a working decoder. A live 1024-by-1024 image render completed through Signal's submission/history/output path. The image was visually inspected for composition and artifacts. This validates that particular image path; it does not establish video throughput or comparative image quality across models.

The generated validation image used a neutral workstation, a lemon, warm lighting, and deliberate space for separately typeset copy. It avoided invented application UI and unreadable generated branding. Its purpose was to test an illustrative asset, not to document a product feature. Product screenshots remain a separate source of evidence.

FLUX.1 Schnell's publisher identifies the model as Apache-2.0. [14] Krea's community license has materially different commercial conditions, including a company-wide trailing annual revenue threshold below $1 million and separate licensing above that threshold. [17] These models should not be presented as having interchangeable commercial terms. Signal links the relevant license next to each recipe rather than silently accepting terms or distributing model weights.

Video requires additional discipline. A short clip should depict one coherent shot; a full promotional video also needs editing, text, audio, pacing, and a final call to action. The campaign studio generates shot-by-shot scripts, while ComfyUI renders an individual configured workflow. CineGen itself can be used to assemble and review the final production. Signal does not call a script or queued render a finished published video.

Scheduling, distribution, and ownership

Agent work and distribution are separate state transitions. A task can finish after producing a useful artifact; an asset can be approved without having been published. Calendar entries are dated distribution reminders, and a publication record requires an actual URL. This distinction prevents dashboards from turning planned work into fictitious reach.

The existing community reply pipeline retains moderator-rule checks, affiliation disclosure, owner review, and publication-specific limits. Campaign studio is a general materials workspace and does not bypass that pipeline. Email drafts are prepared for a consented audience and the owner's sending provider. Platform credentials, unsubscribe controls, and current platform policies remain properties of the actual delivery integration.

The blog integration is a separate configured website path. Its existence should not be taken as proof that every business site, key, or live publication route is operational. The app's broader production workflows remain useful with manual distribution and exports, while service-specific automation can be connected through explicit, authenticated adapters.

Schedules also need an honest runtime boundary. Desktop workers run while Signal is open. CLI workers can run the same task engine in a persistent local process. A powered-off machine cannot run local inference or publish scheduled material. No background service, cloud scheduler, or unattended operating-system startup is implied by saving a date in the calendar.

Measurement and financial learning

The measurement chain is campaign, destination, visits, leads, activation, sale, refund, and retention. Different businesses define activation differently, so Signal records the source and observation rather than assuming every signup is valuable. For Lemonade, a completed local document workflow may be a better activation event than installation. For CineGen, a completed edit or export may be more useful than account creation.

UTM parameters carry campaign metadata, but they do not themselves collect analytics or prove a purchase. Google documents manual campaign parameters and the interaction with automatic advertising tags.[21] The destination website or checkout must retain the necessary information, or the owner must provide a voluntary source answer and reconciliation. For a privacy-positioned desktop product, silently adding behavioral tracking would conflict with the offer being marketed.

Outcome imports are incremental, source-backed records with unique references. Repeating an identical import is harmless; reusing a reference with different values is rejected. This protects against duplicate revenue. Overlapping cumulative analytics exports still require the owner to define disjoint windows, because no local application can infer whether two differently named external reports cover the same visitors.

Sales counts represent transactions unless the source explicitly provides distinct customers. Refund-adjusted revenue is not contribution margin. The spending envelope therefore remains separate: it starts at zero and increases only through recorded settled surplus after costs and obligations. This prevents an

attractive gross-revenue chart from authorizing unaffordable marketing spend.

Industry email benchmarks can contextualize a campaign, but they should not replace the business's own activation and revenue measurements. Mailchimp describes benchmarking as a comparison against similar peers and ties availability to its own analytics program.[22] A universal open-rate target would be an especially weak control for a product with no established audience or purchase baseline.

Recommended operating cadence

The first cycle should produce evidence and a usable invitation, not a large content backlog. Confirm the business profile, choose a narrow buyer problem, prepare a small demonstration, and distribute it in one appropriate place. Record the actual conversation and the next action. Use the results to refine the offer before broadening production.

A practical subsequent cycle is to review unanswered customer questions, inspect the campaign outcomes, identify the most consequential uncertainty, and produce one new asset or improvement request. An agent should reuse the existing campaign when possible rather than generating another duplicate plan. More output is justified when a working message needs adaptation to another relevant destination.

At the weekly review, inspect the full funnel, actual costs, repeated objections, and retention signals. Continue promising tests with adequate evidence, revise tests that reached the wrong audience, and stop tests that consumed their preset resources without useful learning. Move a channel into focused execution only after observed results support it. Keep pricing, partner commission terms, and renewal policies under owner control.

Validation and remaining external prerequisites

The baseline checkout passed TypeScript and 494 tests. Added validation covers provider routing, private-key handling at the transport boundary, daily request limits, campaign recovery, cross-product scoping, review transitions, outcome idempotency, cancellation, and channel coverage. Built Electron checks exercised the workspace, approval persistence, settings persistence, creative setup, and both desktop and narrow layouts. Final command results are recorded in the delivery notes.

A real local model produced persisted social and image-brief assets for both Lemonade and CineGen. Inspection exposed invented product-screen directions and overly absolute privacy copy. The producer now repairs those patterns, requires affiliation in promotional copy, and preserves owner review for claims that deterministic checks cannot establish. A real Krea 2 render produced a saved 1024px image through ComfyUI and was visually inspected. These are bounded live validations, not a claim that all possible prompts or every model will work. GPT requires the owner's API key and account/model access for a live generation test. Wan video requires its own installed weights and a live rendering validation before throughput can be claimed.

Automatic revenue ingestion and distribution depend on the corresponding business systems. Manual evidence entry, idempotent JSON imports, campaign export, and publication recording are available without those credentials. No messages were sent, ad budgets spent, or campaign material published as part of

application validation.

Sources

1. Julian Shapiro. Customer Acquisition Channels. Undated practitioner guide; accessed September 13, 2026.

2. Y Combinator. New YC video series: First Five. July 13, 2022; qualitative founder histories.

3. Lemonade / AIon X. Lemonade product website. Current public positioning and offer, accessed September 13, 2026.

4. CineGen. CineGen product website. Current public positioning and free offer, accessed September 13, 2026.

5. Julian Shapiro. Product-Led Acquisition. Undated practitioner guide; accessed September 13, 2026.

6. Google Search Central. Creating helpful, reliable, people-first content. Living documentation.

7. Google Search Central. Guidance on generative AI content. Living documentation.

8. Ollama. Generate a chat message. Native chat, formatting, runtime options, and keep-alive contract.

9. Ollama. Pull a model. Model download API.

10. Ollama. FAQ. Context configuration, server controls, and cloud-feature configuration.

11. OpenAI. Structured model outputs. JSON mode and strict schema support.

12. Electron. safeStorage. OS encryption semantics and Linux fallback limitations. Implementation uses the synchronous API supported by the installed Electron version.

13. Comfy Org. ComfyUI server routes. Submission, queue, history, and output routes.

14. Black Forest Labs. FLUX.1 Schnell model card. Model characteristics and license.

15. Comfy Org. FLUX.1 text-to-image workflow. Required components and workflow variants.

16. Comfy Org. Krea 2 workflow. Model variants, sampling, and native workflow documentation.

17. Krea AI. Krea 2 Community License. Commercial conditions and distribution requirements; accessed September 13, 2026.

18. Comfy Org. Wan 2.2 native workflows. 5B offloading guidance and required model components.

19. Wan AI. Wan2.2-TI2V-5B model card. Model capabilities and Apache-2.0 identification.

20. LTX. Text-to-video workflow. Current synchronized audio/video tutorial; separate integration candidate.

21. Google Analytics. Traffic-source dimensions, manual tagging, and auto-tagging. Campaign attribution metadata and tagging interactions.

22. Mailchimp. About campaign benchmarking. Peer comparison and service-specific requirements.